

Gruppe 5 - Übung 5

Von: Ivo Janowitz, Nguyen Anh Quang, Tillman Rossa, Roman Seiler, Alexander Uhl

1. Timespow

Aufgabe 1

- Schreiben Sie eine Lösung für den kurzen Test timespow.s aus der letzten Veranstaltung. Der Code befindet sich im Git Repository sysprog

```
# Systemnahe Programmierung - Testen Sie Ihre Kenntnisse!
# H. Hoegl, 2012-11-08

# timespow(3, 2) + timespow(2, 3)

.section .data
.section .text
.globl _start

_start:
    pushl $1          # b
    pushl $7          # x
    call timespow2
    addl $8, %esp
    pushl %eax
    pushl $0          # b
    pushl $2          # x
    call timespow2
    addl $8, %esp
    popl %ebx
    addl %eax, %ebx
    movl $1, %eax
    int $0x80

# timespow2(x, b)
# return x * 2^b
# Trick: x * 2^b = shift argument x left by b bits
#                shll %cl, %ebx (shift ebx left by cl bits)
.type timespow2, @function
timespow2:
    pushl %ebp        # 1 Prolog
    movl %esp, %ebp   # 2 Prolog
    movl 12(%ebp), %ebx # 3 Argument holen
    movb 8(%ebp), %cl  # 4 Argument holen
    shll %cl, %ebx     # 5 Schieben
    movl %ebx, %eax    # 6 Ergebnis ablegen
    movl %ebp, %esp    # 7 Epilog
    popl %ebp         # 8 Epilog
    ret               # 9 Zurueckkehren
```

2. Selbst Übung

Aufgabe 2

- Vollziehen Sie das im Kapitel 5 (Bartlett) beschriebene Programm mit dem Debugger gdb nach

3. ToUpper in anderen Sprachen

Aufgabe 3

- Formulieren Sie das Programm aus Kapitel 5 in den Sprachen

ToUpper in C

C Code to be added

ToUpper in Java

Java Code to be addedd

ToUpper in Python

Python Code to be added

4. Aufgaben von Kapitel 5 PGU

Know the concepts

- Describe the lifecycle of a file descriptor.

A file descriptor is created when we open a file until we close it.

- What are the standard file descriptors and what are they used for?

File Descriptor 0, 1 and 2 are the standard one. 0 is the STDIN, 1 is the STDOUT and 2 is the STDERR. Where STDIN represent the default keyboard, STDOUT the screen and STDERR is the standard error

- What is a buffer?

A Buffer is a place to store a big junk of data for example from the read of a file

- What is the difference between the .data section and the .bss section?

The .data section is static and the .bss is dynamic for one example. Also the .data section stores the initialized data from the source code where the .bss holds temporary new data create in the runtime

- What are the system calls related to reading and writing files?

```
movl $5, %eax    #0pen
int 0x80
```

```

movl $4, %eax    #Write
int 0x80

movl $3, %eax    #Read
int 0x80

movl $6, %eax    #Close
int 0x80

```

Use the concepts

- Modify the toupper program so that it reads from STDIN and writes to STDOUT instead of using the files on the command-line.

```

# -*- indent-tabs-mode: nil -*- (for Emacs)

# PURPOSE: This program converts an input file
# to an output file with all letters
# converted to uppercase.
#
# PROCESSING:
# 1) Open the input file
# 2) Open the output file
# 4) While we're not at the end of the input file
#    a) read part of file into our memory buffer
#    b) go through each byte of memory
#       if the byte is a lower-case letter,
#       convert it to uppercase
#    c) write the memory buffer to output file

.section .data

#####CONSTANTS#####
# system call numbers
.equ SYS_OPEN, 5
.equ SYS_WRITE, 4
.equ SYS_READ, 3
.equ SYS_CLOSE, 6
.equ SYS_EXIT, 1

# options for open (look at
# /usr/include/asm/fcntl.h for
# various values. You can combine them
# by adding them or ORing them)
# This is discussed at greater length
# in "Counting Like a Computer"
.equ O_RDONLY, 0
.equ O_CREAT_WRONLY_TRUNC, 03101

# standard file descriptors
.equ STDIN, 0
.equ STDOUT, 1
.equ STDERR, 2

# system call interrupt
.equ LINUX_SYSCALL, 0x80
.equ END_OF_FILE, 0 #This is the return value

# of read which means we've
# hit the end of the file
.equ NUMBER_ARGUMENTS, 2

.section .bss

# Buffer - this is where the data is loaded into
# from the data file and written from
# into the output file. This should
# never exceed 16,000 for various
# reasons.
.equ BUFFER_SIZE, 500
.lcomm BUFFER_DATA, BUFFER_SIZE

```

```

.section .text

# STACK POSITIONS
.equ ST_SIZE_RESERVE, 8
.equ ST_FD_IN, -4
.equ ST_FD_OUT, -8
.equ ST_ARGC, 0          # Number of arguments
.equ ST_ARGV_0, 4        # Name of program
.equ ST_ARGV_1, 8        # Input file name
.equ ST_ARGV_2, 12       # Output file name
.globl _start

_start:
### INITIALIZE PROGRAM ###
# save the stack pointer

movl %esp, %ebp

# Allocate space for our file descriptors
# on the stack
subl $ST_SIZE_RESERVE, %esp

store_fd_in:
# save the given file descriptor
movl $0, ST_FD_IN(%ebp)

store_fd_out:
# store the file descriptor here
movl $1, ST_FD_OUT(%ebp)

                                # BEGIN MAIN LOOP
read_loop_begin:
                                # READ IN A BLOCK FROM THE INPUT FILE

movl $SYS_READ, %eax
movl ST_FD_IN(%ebp), %ebx      # get the input file descriptor
movl $BUFFER_DATA, %ecx       # the location to read into
movl $BUFFER_SIZE, %edx       # the size of the buffer
int $LINUX_SYSCALL            # Size of buffer read is returned in %eax
cmpl $END_OF_FILE, %eax       # check for end of file marker
                                # if found or on error, go to the end
jle end_loop

continue_read_loop:
### CONVERT THE BLOCK TO UPPER CASE ###
pushl $BUFFER_DATA            # location of buffer
pushl %eax                     # size of the buffer
call convert_to_upper
popl %eax                      # get the size back
addl $4, %esp                  # restore %esp
### WRITE THE BLOCK OUT TO THE OUTPUT FILE ###
movl %eax, %edx                # size of the buffer
movl $SYS_WRITE, %eax
movl ST_FD_OUT(%ebp), %ebx     # file to use
movl $BUFFER_DATA, %ecx       # location of the buffer
int $LINUX_SYSCALL
jmp read_loop_begin

end_loop:
                                ###CLOSE THE FILES###
                                # NOTE - we don't need to do error checking
                                # on these, because error conditions
                                # don't signify anything special here

movl $SYS_CLOSE, %eax
movl ST_FD_OUT(%ebp), %ebx
int $LINUX_SYSCALL
movl $SYS_CLOSE, %eax
movl ST_FD_IN(%ebp), %ebx
int $LINUX_SYSCALL

                                ### EXIT ###
movl $SYS_EXIT, %eax
movl $0, %ebx
int $LINUX_SYSCALL

#PURPOSE: This function actually does the
# conversion to upper case for a block
#
#INPUT: The first parameter is the location

```

```

# of the block of memory to convert
# The second parameter is the length of
# that buffer
#
#OUTPUT: This function overwrites the current
# buffer with the upper-casified version.
#
#VARIABLES:
# %eax - beginning of buffer
# %ebx - length of buffer
# %edi - current buffer offset
# %cl - current byte being examined
# (first part of %ecx)

    ### CONSTANTS ##
    .equ LOWERCASE_A, 'a'      # The lower boundary of our search
    .equ LOWERCASE_Z, 'z'      # The upper boundary of our search
    .equ UPPER_CONVERSION, 'A' - 'a' # Conversion between upper and lower case
    ### STACK STUFF ###
    .equ ST_BUFFER_LEN, 8      # Length of buffer
    .equ ST_BUFFER, 12         # actual buffer

convert_to_upper:
    pushl %ebp
    movl %esp, %ebp

    ### SET UP VARIABLES ###
    movl ST_BUFFER(%ebp), %eax
    movl ST_BUFFER_LEN(%ebp), %ebx
    movl $0, %edi

    # if a buffer with zero length was given
    # to us, just leave
    cmpl $0, %ebx
    je end_convert_loop

convert_loop:
    # get the current byte
    movb (%eax,%edi,1), %cl
    # go to the next byte unless it is between
    # 'a' and 'z'
    cmpb $LOWERCASE_A, %cl
    jl next_byte
    cmpb $LOWERCASE_Z, %cl
    jg next_byte

    # otherwise convert the byte to uppercase
    addb $UPPER_CONVERSION, %cl
    # and store it back
    movb %cl, (%eax,%edi,1)

next_byte:
    incl %edi #next byte
    cmpl %edi, %ebx #continue unless
    # we've reached the end
    jne convert_loop
end_convert_loop:

    # no return value, just leave
    movl %ebp, %esp
    popl %ebp
    ret

# vim: expandtab sw=8 sts=8

```

ToUpper Buffer Version

- Change the size of the buffer.

```

# -*- indent-tabs-mode: nil -*- (for Emacs)

# PURPOSE: This program converts an input file
# to an output file with all letters
# converted to uppercase.
#
# PROCESSING:
# 1) Open the input file

```

```

# 2) Open the output file
# 4) While we're not at the end of the input file
#     a) read part of file into our memory buffer
#     b) go through each byte of memory
#         if the byte is a lower-case letter,
#         convert it to uppercase
#     c) write the memory buffer to output file

.section .data

#####CONSTANTS#####
# system call numbers
.equ SYS_OPEN, 5
.equ SYS_WRITE, 4
.equ SYS_READ, 3
.equ SYS_CLOSE, 6
.equ SYS_EXIT, 1

# options for open (look at
# /usr/include/asm/fcntl.h for
# various values. You can combine them
# by adding them or ORing them)
# This is discussed at greater length
# in "Counting Like a Computer"
.equ O_RDONLY, 0
.equ O_CREAT_WRONLY_TRUNC, 03101

# standard file descriptors
.equ STDIN, 0
.equ STDOUT, 1
.equ STDERR, 2

# system call interrupt
.equ LINUX_SYSCALL, 0x80
.equ END_OF_FILE, 0 #This is the return value

# of read which means we've
# hit the end of the file
.equ NUMBER_ARGUMENTS, 2

.section .bss

# Buffer - this is where the data is loaded into
# from the data file and written from
# into the output file. This should
# never exceed 16,000 for various
# reasons.
.equ BUFFER_SIZE, 20

#.equ BUFFER_SIZE, 100
#.equ BUFFER_SIZE, 500
#.equ BUFFER_SIZE, 100000

.lcomm BUFFER_DATA, BUFFER_SIZE

.section .text

# STACK POSITIONS
.equ ST_SIZE_RESERVE, 8
.equ ST_FD_IN, -4
.equ ST_FD_OUT, -8
.equ ST_ARGC, 0           # Number of arguments
.equ ST_ARGV_0, 4         # Name of program
.equ ST_ARGV_1, 8         # Input file name
.equ ST_ARGV_2, 12        # Output file name
.globl _start

_start:
### INITIALIZE PROGRAM ###
# save the stack pointer

movl %esp, %ebp

# Allocate space for our file descriptors
# on the stack
subl $ST_SIZE_RESERVE, %esp

```

```

open_files:
open_fd_in:
    ### OPEN INPUT FILE ###
    # open syscall
    movl $SYS_OPEN, %eax
    # input filename into %ebx
    movl ST_ARGV_1(%ebp), %ebx
    # read-only flag
    movl $O_RDONLY, %ecx
    # this doesn't really matter for reading
    movl $0666, %edx
    # call Linux
    int $LINUX_SYSCALL
store_fd_in:
    # save the given file descriptor
    movl %eax, ST_FD_IN(%ebp)

open_fd_out:
    ### OPEN OUTPUT FILE ###
    # open the file
    movl $SYS_OPEN, %eax
    # output filename into %ebx
    movl ST_ARGV_2(%ebp), %ebx
    # flags for writing to the file
    movl $O_CREAT_WRONLY_TRUNC, %ecx
    # mode for new file (if it's created)
    movl $0666, %edx
    # call Linux
    int $LINUX_SYSCALL

store_fd_out:
    # store the file descriptor here
    movl %eax, ST_FD_OUT(%ebp)
    # BEGIN MAIN LOOP

read_loop_begin:
    # READ IN A BLOCK FROM THE INPUT FILE
    movl $SYS_READ, %eax
    movl ST_FD_IN(%ebp), %ebx # get the input file descriptor
    movl $BUFFER_DATA, %ecx # the location to read into
    movl $BUFFER_SIZE, %edx # the size of the buffer
    int $LINUX_SYSCALL # Size of buffer read is returned in %eax
    cmpl $END_OF_FILE, %eax # check for end of file marker
    # if found or on error, go to the end
    jle end_loop
continue_read_loop:
    ### CONVERT THE BLOCK TO UPPER CASE ###
    pushl $BUFFER_DATA # location of buffer
    pushl %eax # size of the buffer
    call convert_to_upper
    popl %eax # get the size back
    addl $4, %esp # restore %esp
    ### WRITE THE BLOCK OUT TO THE OUTPUT FILE ###
    movl %eax, %edx # size of the buffer
    movl $SYS_WRITE, %eax
    movl ST_FD_OUT(%ebp), %ebx # file to use
    movl $BUFFER_DATA, %ecx # location of the buffer
    int $LINUX_SYSCALL
    jmp read_loop_begin

end_loop:
    ###CLOSE THE FILES###
    # NOTE - we don't need to do error checking
    # on these, because error conditions
    # don't signify anything special here

    movl $SYS_CLOSE, %eax
    movl ST_FD_OUT(%ebp), %ebx
    int $LINUX_SYSCALL
    movl $SYS_CLOSE, %eax
    movl ST_FD_IN(%ebp), %ebx
    int $LINUX_SYSCALL

    ### EXIT ###
    movl $SYS_EXIT, %eax
    movl $0, %ebx
    int $LINUX_SYSCALL

```

#PURPOSE: This function actually does the
conversion to upper case for a block

```

#
#INPUT: The first parameter is the location
# of the block of memory to convert
# The second parameter is the length of
# that buffer
#
#OUTPUT: This function overwrites the current
# buffer with the upper-casified version.
#
#VARIABLES:
# %eax - beginning of buffer
# %ebx - length of buffer
# %edi - current buffer offset
# %cl - current byte being examined
# (first part of %ecx)

    ### CONSTANTS ##
    .equ LOWERCASE_A, 'a'          # The lower boundary of our search
    .equ LOWERCASE_Z, 'z'          # The upper boundary of our search
    .equ UPPER_CONVERSION, 'A' - 'a' # Conversion between upper and lower case
    ### STACK STUFF ###
    .equ ST_BUFFER_LEN, 8          # Length of buffer
    .equ ST_BUFFER, 12             # actual buffer
convert_to_upper:
    pushl %ebp
    movl %esp, %ebp
    ### SET UP VARIABLES ###
    movl ST_BUFFER(%ebp), %eax
    movl ST_BUFFER_LEN(%ebp), %ebx
    movl $0, %edi

    # if a buffer with zero length was given
    # to us, just leave
    cmpl $0, %ebx
    je end_convert_loop
convert_loop:
    # get the current byte
    movb (%eax,%edi,1), %cl
    # go to the next byte unless it is between
    # 'a' and 'z'
    cmpb $LOWERCASE_A, %cl
    jl next_byte
    cmpb $LOWERCASE_Z, %cl
    jg next_byte
    # otherwise convert the byte to uppercase
    addb $UPPER_CONVERSION, %cl
    # and store it back
    movb %cl, (%eax,%edi,1)

next_byte:
    incl %edi #next byte
    cmpl %edi, %ebx #continue unless
    # we've reached the end
    jne convert_loop
end_convert_loop:
    # no return value, just leave
    movl %ebp, %esp
    popl %ebp
    ret

# vim: expandtab sw=8 sts=8

```

ToUpper BSS Version

- Rewrite the program so that it uses storage in the .bss section rather than the stack to store the file descriptors.

```

# -*- indent-tabs-mode: nil -*- (for Emacs)

# PURPOSE: This program converts an input file

```



```

# to an output file with all letters
# converted to uppercase.
#
# PROCESSING:
# 1) Open the input file
# 2) Open the output file
# 4) While we're not at the end of the input file
#   a) read part of file into our memory buffer
#   b) go through each byte of memory
#       if the byte is a lower-case letter,
#       convert it to uppercase
#   c) write the memory buffer to output file

.section .data

#####CONSTANTS#####
# system call numbers
.equ SYS_OPEN, 5
.equ SYS_WRITE, 4
.equ SYS_READ, 3
.equ SYS_CLOSE, 6
.equ SYS_EXIT, 1

# options for open (look at
# /usr/include/asm/fcntl.h for
# various values. You can combine them
# by adding them or ORing them)
# This is discussed at greater length
# in "Counting Like a Computer"
.equ O_RDONLY, 0
.equ O_CREAT_WRONLY_TRUNC, 03101

# standard file descriptors
.equ STDIN, 0
.equ STDOUT, 1
.equ STDERR, 2

# system call interrupt
.equ LINUX_SYSCALL, 0x80
.equ END_OF_FILE, 0 #This is the return value

# of read which means we've
# hit the end of the file
.equ NUMBER_ARGUMENTS, 2

.section .bss

# Buffer - this is where the data is loaded into
# from the data file and written from
# into the output file. This should
# never exceed 16,000 for various
# reasons.
.equ BUFFER_SIZE, 500
.lcomm BUFFER_DATA, BUFFER_SIZE
.lcomm FILEIN, 4
.lcomm FILEOUT, 4

.section .text

# STACK POSITIONS
.equ ST_SIZE_RESERVE, 8
.equ ST_FD_IN, -4
.equ ST_FD_OUT, -8
.equ ST_ARGC, 0          # Number of arguments
.equ ST_ARGV_0, 4        # Name of program
.equ ST_ARGV_1, 8        # Input file name
.equ ST_ARGV_2, 12       # Output file name
.globl _start

_start:
### INITIALIZE PROGRAM ###
# save the stack pointer

movl %esp, %ebp

# Allocate space for our file descriptors
# on the stack

```

```

        subl $ST_SIZE_RESERVE, %esp

open_files:
open_fd_in:
        ### OPEN INPUT FILE ###
        # open syscall
        movl $SYS_OPEN, %eax
        # input filename into %ebx
        movl ST_ARGV_1(%ebp), %ebx
        # read-only flag
        movl $O_RDONLY, %ecx
        # this doesn't really matter for reading
        movl $0666, %edx
        # call Linux
        int $LINUX_SYSCALL
store_fd_in:
        # save the given file descriptor
        movl %eax, FILEIN

open_fd_out:
        ### OPEN OUTPUT FILE ###
        # open the file
        movl $SYS_OPEN, %eax
        # output filename into %ebx
        movl ST_ARGV_2(%ebp), %ebx
        # flags for writing to the file
        movl $O_CREAT|W_WRONLY|TRUNC, %ecx
        # mode for new file (if it's created)
        movl $0666, %edx
        # call Linux
        int $LINUX_SYSCALL

store_fd_out:
        # store the file descriptor here
        movl %eax, FILEOUT

        # BEGIN MAIN LOOP
read_loop_begin:
        # READ IN A BLOCK FROM THE INPUT FILE
        movl $SYS_READ, %eax
        movl FILEIN, %ebx
        # get the input file descriptor
        movl $BUFFER_DATA, %ecx
        # the location to read into
        movl $BUFFER_SIZE, %edx
        # the size of the buffer
        int $LINUX_SYSCALL
        # Size of buffer read is returned in %eax
        cmpl $END_OF_FILE, %eax
        # check for end of file marker
        # if found or on error, go to the end
        jle end_loop
continue_read_loop:
        ### CONVERT THE BLOCK TO UPPER CASE ###
        pushl $BUFFER_DATA
        # location of buffer
        pushl %eax
        # size of the buffer
        call convert_to_upper
        popl %eax
        # get the size back
        addl $4, %esp
        # restore %esp
        ### WRITE THE BLOCK OUT TO THE OUTPUT FILE ###
        movl %eax, %edx
        # size of the buffer
        movl $SYS_WRITE, %eax
        movl FILEOUT, %ebx
        # file to use
        movl $BUFFER_DATA, %ecx
        # location of the buffer
        int $LINUX_SYSCALL
        jmp read_loop_begin

end_loop:
        ###CLOSE THE FILES###
        # NOTE - we don't need to do error checking
        # on these, because error conditions
        # don't signify anything special here

        movl $SYS_CLOSE, %eax
        movl FILEOUT, %ebx
        int $LINUX_SYSCALL
        movl $SYS_CLOSE, %eax
        movl FILEIN, %ebx
        int $LINUX_SYSCALL

        ### EXIT ###
        movl $SYS_EXIT, %eax
        movl $0, %ebx
        int $LINUX_SYSCALL

```

```

#PURPOSE: This function actually does the
# conversion to upper case for a block
#
#INPUT: The first parameter is the location
#
# of the block of memory to convert
# The second parameter is the length of
# that buffer
#
#OUTPUT: This function overwrites the current
# buffer with the upper-casified version.
#
#VARIABLES:
# %eax - beginning of buffer
# %ebx - length of buffer
# %edi - current buffer offset
# %cl - current byte being examined
# (first part of %ecx)

    ### CONSTANTS ##
    .equ LOWERCASE_A, 'a'      # The lower boundary of our search
    .equ LOWERCASE_Z, 'z'      # The upper boundary of our search
    .equ UPPER_CONVERSION, 'A' - 'a' # Conversion between upper and lower case
    ### STACK STUFF ###
    .equ ST_BUFFER_LEN, 8      # Length of buffer
    .equ ST_BUFFER, 12         # actual buffer

convert_to_upper:
    pushl %ebp
    movl %esp, %ebp
    ### SET UP VARIABLES ###
    movl ST_BUFFER(%ebp), %eax
    movl ST_BUFFER_LEN(%ebp), %ebx
    movl $0, %edi

    # if a buffer with zero length was given
    # to us, just leave
    cmpl $0, %ebx
    je end_convert_loop
convert_loop:
    # get the current byte
    movb (%eax,%edi,1), %cl
    # go to the next byte unless it is between
    # 'a' and 'z'
    cmpb $LOWERCASE_A, %cl
    jl next_byte
    cmpb $LOWERCASE_Z, %cl
    jg next_byte
    # otherwise convert the byte to uppercase
    addb $UPPER_CONVERSION, %cl
    # and store it back
    movb %cl, (%eax,%edi,1)

next_byte:
    incl %edi #next byte
    cmpl %edi, %ebx #continue unless
    # we've reached the end
    jne convert_loop
end_convert_loop:
    # no return value, just leave
    movl %ebp, %esp
    popl %ebp
    ret

# vim: expandtab sw=8 sts=8

```

HeyNow

- Write a program that will create a file called [heyNow.txt](#) and write the words “Hey diddle diddle!” into it.

```

# -*- indent-tabs-mode: nil -*- (for Emacs)

# PURPOSE:

        .section .data

        # File name and contents
file_name:
        .ascii "heynow.txt\0"
file_contents:
        .ascii "Hey diddle diddle!\0"

#####CONSTANTS#####
# system call numbers
.equ SYS_OPEN, 5
.equ SYS_WRITE, 4
.equ SYS_READ, 3
.equ SYS_CLOSE, 6
.equ SYS_EXIT, 1

# options for open (look at
# /usr/include/asm/fcntl.h for
# various values. You can combine them
# by adding them or ORing them)
# This is discussed at greater length
# in "Counting Like a Computer"
.equ O_RDONLY, 0
.equ O_CREAT_WRONLY_TRUNC, 03101

# standard file descriptors
.equ STDIN, 0
.equ STDOUT, 1
.equ STDERR, 2

# system call interrupt
.equ LINUX_SYSCALL, 0x80
.equ END_OF_FILE, 0 #This is the return value

# of read which means we've
# hit the end of the file
.equ NUMBER_ARGUMENTS, 2

.section .bss

# Buffer - this is where the data is loaded into
# from the data file and written from
# into the output file. This should
# never exceed 16,000 for various
# reasons.
.equ BUFFER_SIZE, 500
.lcomm BUFFER_DATA, BUFFER_SIZE

.section .text

# STACK POSITIONS
.equ ST_SIZE_RESERVE, 8
.equ ST_FD_OUT, -4
.globl _start

_start:
    ### INITIALIZE PROGRAM ###
    # save the stack pointer

    movl %esp, %ebp

    # Allocate space for our file descriptors
    # on the stack
    subl $ST_SIZE_RESERVE, %esp

open_files:
open_fd_out:
                                ### OPEN OUTPUT FILE ###
                                # open the file

    movl $SYS_OPEN, %eax
                                # output filename into %ebx
    movl $file_name, %ebx
                                # flags for writing to the file
    movl $O_CREAT_WRONLY_TRUNC, %ecx

```

```

                                # mode for new file (if it's created)
                                # call Linux
                                # store the file descriptor here
                                # text to write
                                # counter
                                # if zero byte, string is done
                                # increment
                                # if zero byte, string is done
                                # size of the buffer
                                # file to use
                                # location of the buffer
                                ###CLOSE THE FILES###
                                # NOTE - we don't need to do error checking
                                # on these, because error conditions
                                # don't signify anything special here
                                ### EXIT ###
# vim: expandtab sw=8 sts=8

```